

Group 10

COMPARISON OF NEURAL NETWORK ARCHITECTURES FOR ART GENRE CLASSIFICATION USING IMAGE METADATA

**HOON HAN
JOONHYOUNG LEE
STEPHEN YEAP WELLER**

TABLE OF CONTENTS

- 1. INTRODUCTION**
- 2. DATASET OVERVIEW**
- 3. DATA PREPROCESSING**
- 4. MODEL'S ARCHITECTURE**
- 5. EVALUATIONS & RESULTS**
- 6. LIMITATIONS &
POTENTIAL IMPROVEMENTS**

Introduction

Project Objective

- Classify paintings into artistic genres using machine learning.
- Focus on stylistic patterns over artist identification.

Our Approach

- **Simple Neural Network (SNN)**
 - Baseline model with straightforward architecture.
- **Boosted Neural Network (BNN)**
 - Ensemble method to correct misclassification.
- **SNN with Adam Optimizer**
 - Adaptive learning rate for improved convergence.

Significance

- **Art Education:** Tools for learning and exploring art styles.
- **Digital Curation:** Efficient organization of art collections for professionals and enthusiasts.

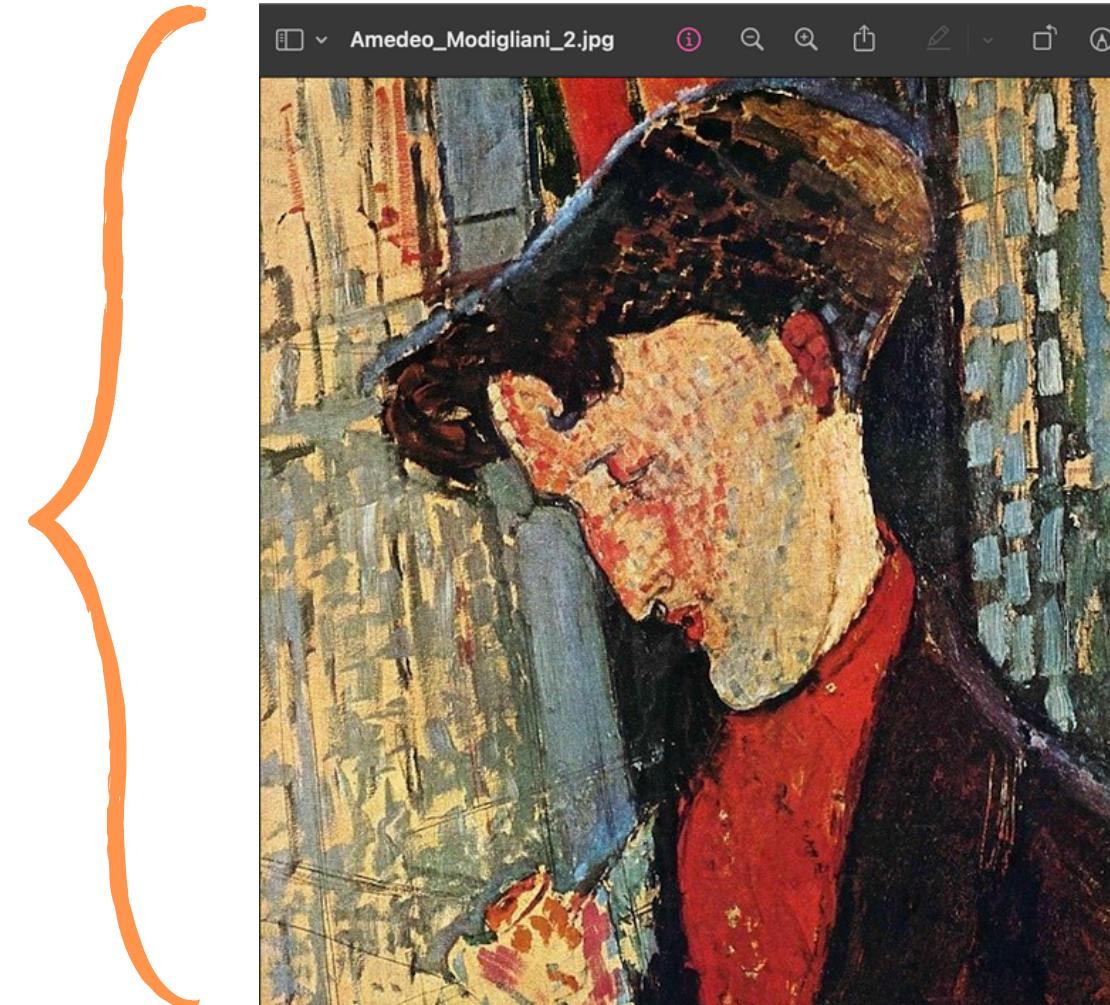


Dataset Overview

artists

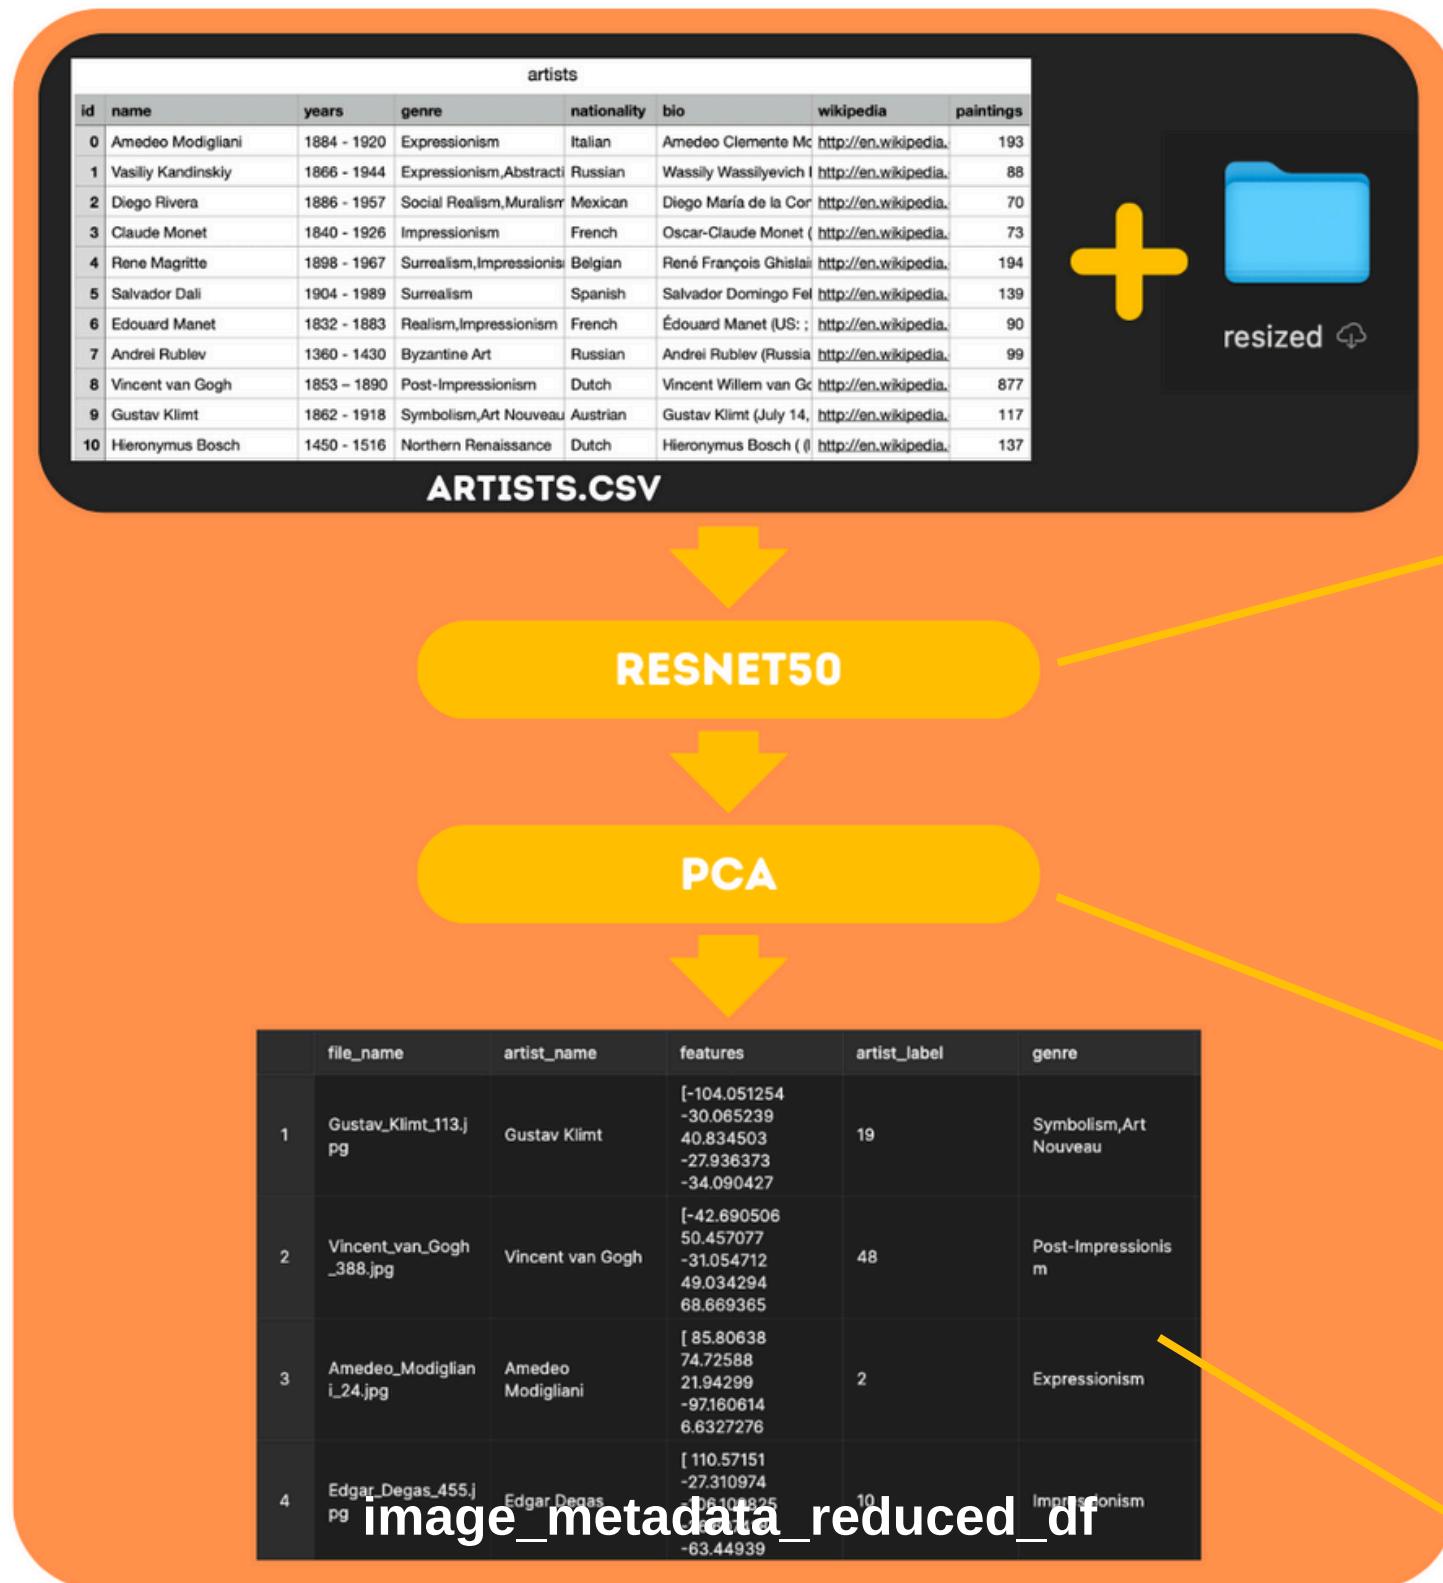
id	name	years	genre	nationality	bio	wikipedia	paintings
0	Amedeo Modigliani	1884 - 1920	Expressionism	Italian	Amedeo Clemente Modigliani (Italian pronunciation: [amedeo mɔdiljaːni]; 12 September 1884 – 24 January 1920) was an Italian painter and sculptor who lived most of his life in France. He painted portraits, figures, and nudes, and made sculptures and drawings. His style was influenced by Cubism and Fauvism.	http://en.wikipedia.org/wiki/Amedeo_Modigliani	193
1	Vasily Kandinsky	1866 - 1944	Expressionism, Abstractionism	Russian	Vassily Wassilyevich Kandinsky (Russian: Васи́лий Васи́льевич Кандинский; 4 December 1863 – 13 December 1944) was a Russian painter and art theorist. He is regarded as one of the most important figures in modern art.	http://en.wikipedia.org/wiki/Vassily_Kandinsky	88
2	Diego Rivera	1886 - 1957	Social Realism, Muralism	Mexican	Diego María de la Concepción Juan Nepomuceno Estanislao de la Rivera y Barrientos Acosta y Rodríguez (February 8, 1886 – November 24, 1957) was a Mexican painter, known for his political murals.	http://en.wikipedia.org/wiki/Diego_Rivera	70
3	Claude Monet	1840 - 1926	Impressionism	French	Oscar-Claude Monet (French: [klod mɔnɛ]; 14 November 1840 – 5 December 1926) was a French painter. He is best known for his series of landscapes, particularly those of the River Seine at Giverny.	http://en.wikipedia.org/wiki/Claude_Monet	73
4	Rene Magritte	1898 - 1967	Surrealism, Impressionism	Belgian	René François Ghislain Magritte (French: [ʁəne mɑɡʁit]; 21 November 1898 – 15 August 1967) was a Belgian surrealist painter.	http://en.wikipedia.org/wiki/René_Magritte	194
5	Salvador Dali	1904 - 1989	Surrealism	Spanish	Salvador Domingo Felipe Jacinto Dalí i Domènech (1904–1989) was a Spanish surrealist painter.	http://en.wikipedia.org/wiki/Salvador_Dalí	139
6	Edouard Manet	1832 - 1883	Realism, Impressionism	French	Édouard Manet (US: ; UK: ; French: [edwaʁ mañɛ]) (1832–1883) was a French painter.	http://en.wikipedia.org/wiki/Édouard_Manet	90
7	Andrei Rublev	1360 - 1430	Byzantine Art	Russian	Andrei Rublev (Russian: Андрей Рублёв, IPA: [ɐndrʲej rʊbl̑əf]) (c. 1360 – c. 1430) was a Russian icon painter.	http://en.wikipedia.org/wiki/Andrei_Rublev	99
8	Vincent van Gogh	1853 – 1890	Post-Impressionism	Dutch	Vincent Willem van Gogh (Dutch: [vɪnseːnt ˈvɑŋ ɣɔx]; 30 March 1853 – 29 July 1890) was a Dutch post-impressionist painter.	http://en.wikipedia.org/wiki/Vincent_van_Gogh	877
9	Gustav Klimt	1862 - 1918	Symbolism, Art Nouveau	Austrian	Gustav Klimt (July 14, 1862 – February 6, 1918) was an Austrian symbolist painter.	http://en.wikipedia.org/wiki/Gustav_Klimt	117
10	Hieronymus Bosch	1450 - 1516	Northern Renaissance	Dutch	Hieronymus Bosch ((listen); Dutch: [hijeːrɔmʏs ˈbɔs]) (c. 1450 – 9 August 1516) was a Dutch Northern Renaissance painter.	http://en.wikipedia.org/wiki/Hieronymus_Bosch	137

Artists.csv: 50 Artists, genre labels



resized folder: 8355 image files of paintings

Data Preprocessing



```
# Load the pretrained ResNet50 model (without the top classification layer)
model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
```

Function to extract features from an image

```
def extract_features(image_path):
    img = load_img(image_path, target_size=(224, 224)) # Resize image to 224x224
    img_array = img_to_array(img)
    img_array = preprocess_input(img_array) # Preprocess for ResNet50
    img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
    features = model.predict(img_array)
    return features.flatten() # Flatten the feature map to a 1D vector
```

```
# Step 1: Extract feature matrix from image_metadata_df
```

```
X = np.stack(image_metadata_df['features'].values) # Shape: [n_samples, n_features]
```

Step 2: Apply PCA to reduce dimensionality to 100 components

```
pca = PCA(n_components=100, random_state=42)
X_reduced = pca.fit_transform(X) # Shape: [n_samples, 100]
```

Step 3: Replace the 'features' column with the reduced features

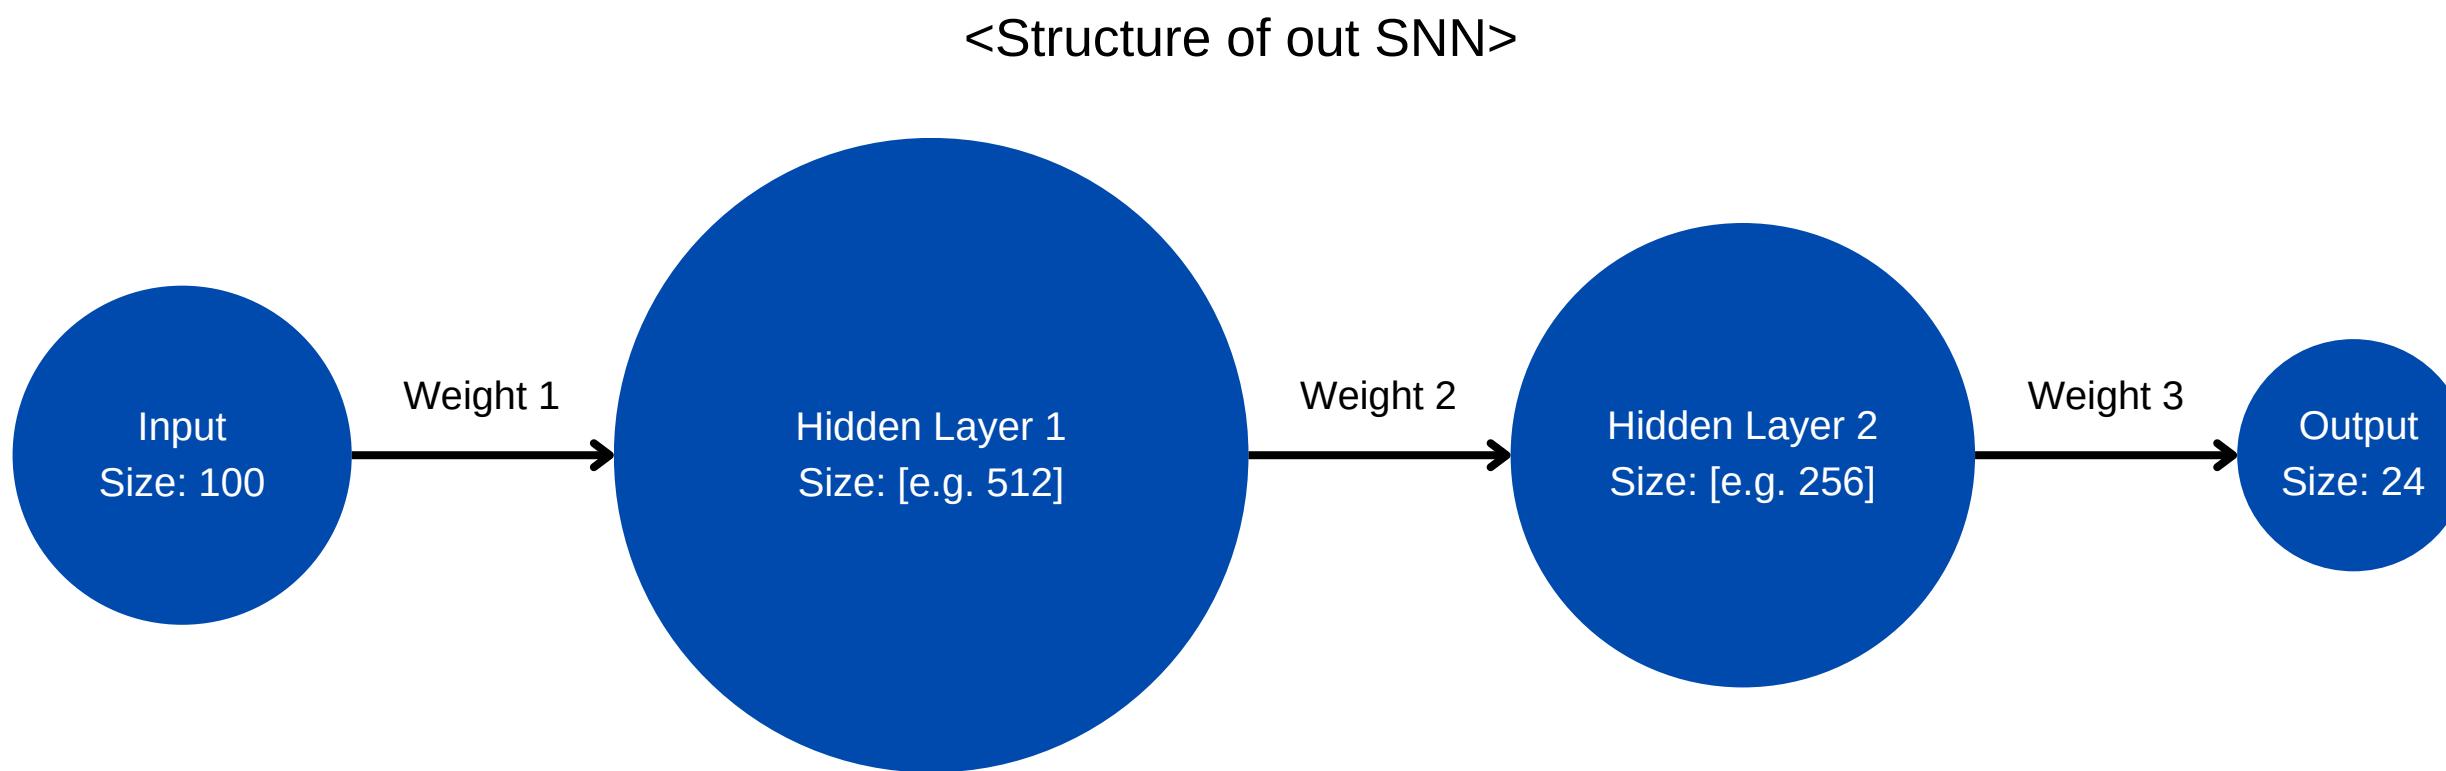
```
image_metadata_df['features'] = list(X_reduced) # Store reduced features as lists
```

One-hot encode the genres

```
mlb = MultiLabelBinarizer()
one_hot_genres = mlb.fit_transform(image_metadata_reduced_df['genre_list'])
```

Models' Architecture

Model 1: Simple Neural Network (SNN)



Why Only Two Layers?

1. Pre-trained ResNet50 & PCA
2. Gradient Vanishing
3. Insufficient data (8,355)
4. Computational cost
5. Imbalanced genre labels

Model 1: Simple Neural Network (SNN)

1) Initialization of Parameters

```
self.weights = {
    "W1": np.random.randn(input_size, hidden_sizes[0]) * 0.01,
    "W2": np.random.randn(hidden_sizes[0], hidden_sizes[1]) * 0.01,
    "W3": np.random.randn(hidden_sizes[1], output_size) * 0.01,
}
self.biases = {
    "b1": np.zeros((1, hidden_sizes[0])),
    "b2": np.zeros((1, hidden_sizes[1])),
    "b3": np.zeros((1, output_size)),
}
```

2) Forward Propagation

```
def forward(self, X):
    # Forward propagation
    self.Z1 = np.dot(X, self.weights["W1"]) + self.biases["b1"]
    self.A1 = self.relu(self.Z1)

    self.Z2 = np.dot(self.A1, self.weights["W2"]) + self.biases["b2"]
    self.A2 = self.relu(self.Z2)

    self.Z3 = np.dot(self.A2, self.weights["W3"]) + self.biases["b3"]
    self.A3 = self.sigmoid(self.Z3) # Sigmoid activation for multi-label classification
```

Theory

$$Z^{[l]} = W^{[l]} A^{[l-1]} + b^{[l]}, \quad A^{[l]} = \text{ReLU}(Z^{[l]}), \quad A^{[L]} = \sigma(Z^{[L]})$$

Model 1: Simple Neural Network (SNN)

3) Loss Computation

```
def binary_cross_entropy_loss(self, y_pred, y_true):
    n_samples = y_true.shape[0]
    return -np.sum(y_true * np.log(y_pred + 1e-10) + (1 - y_true) * np.log(1 - y_pred + 1e-10)) / n_samples
```

Theory

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^n [y_i \log(y_{\text{pred},i}) + (1 - y_i) \log(1 - y_{\text{pred},i})]$$

4) Backward Propagation and Weights & Bias Update

```
def backward(self, X, y_true, y_pred):
    # Backward propagation
    n_samples = y_true.shape[0]

    # Gradients for the output layer
    dZ3 = y_pred - y_true # Binary cross-entropy gradient
    dW3 = np.dot(self.A2.T, dZ3) / n_samples
    db3 = np.sum(dZ3, axis=0, keepdims=True) / n_samples

    # Gradients for the second hidden layer
    dA2 = np.dot(dZ3, self.weights["W3"].T)
    dZ2 = dA2 * self.relu_derivative(self.Z2)
    dW2 = np.dot(self.A1.T, dZ2) / n_samples
    db2 = np.sum(dZ2, axis=0, keepdims=True) / n_samples

    # Update weights and biases
    self.weights["W1"] -= self.learning_rate * dW1
    self.biases["b1"] -= self.learning_rate * db1

    self.weights["W2"] -= self.learning_rate * dW2
    self.biases["b2"] -= self.learning_rate * db2

    self.weights["W3"] -= self.learning_rate * dW3
    self.biases["b3"] -= self.learning_rate * db3
```

Theory

$$\delta^{[L]} = A^{[L]} - Y$$

$$\frac{\partial \mathcal{L}}{\partial W^{[L]}} = A^{[L-1]} \delta^{[L]} \quad \text{and} \quad \frac{\partial \mathcal{L}}{\partial b^{[L]}} = \delta^{[L]}$$

$$\delta^{[l]} = (W^{[l+1]})^T \delta^{[l+1]} \odot \text{ReLU}'(Z^{[l]})$$

$$\frac{\partial \mathcal{L}}{\partial W^{[l]}} = A^{[l-1]} \delta^{[l]} \quad \text{and} \quad \frac{\partial \mathcal{L}}{\partial b^{[l]}} = \delta^{[l]}$$

$$W^{[l]} = W^{[l]} - \eta \frac{\partial \mathcal{L}}{\partial W^{[l]}} \quad \text{and} \quad b^{[l]} = b^{[l]} - \eta \frac{\partial \mathcal{L}}{\partial b^{[l]}}$$

Model 1: Simple Neural Network (SNN)

6) Training Loop

```
def train(self, X, y, epochs=20, batch_size=32):
    n_samples = X.shape[0]

    for epoch in range(epochs):
        # Shuffle the data
        indices = np.arange(n_samples)
        np.random.shuffle(indices)
        X = X[indices]
        y = y[indices]

        # Mini-batch gradient descent
        for i in range(0, n_samples, batch_size):
            X_batch = X[i:i + batch_size]
            y_batch = y[i:i + batch_size]

            # Forward and backward propagation
            y_pred = self.forward(X_batch)
            self.backward(X_batch, y_batch, y_pred)

        # Compute loss
        y_pred_full = self.forward(X)
        loss = self.binary_cross_entropy_loss(y_pred_full, y)
        print(f"Epoch {epoch + 1}/{epochs}, Loss: {loss:.4f}")
```

Theory

$$W^{[l]} = W^{[l]} - \eta \frac{\partial \mathcal{L}}{\partial W^{[l]}} \quad \text{and} \quad b^{[l]} = b^{[l]} - \eta \frac{\partial \mathcal{L}}{\partial b^{[l]}}$$

Theory

- Mini-batching
- Epochs repeat

Model 2: Boosted Neural Network- Ensemble Method (BNN)

Theory

Train multiple neural networks sequentially, with each subsequent model addressing the errors of its predecessors.

1,2,3) Initialization of Parameters, Forward Propagation, Loss Computation

Same with Model 1

4) Backward Propagation and Weight & Bias Updates:

Gradient clipping ensures that gradients do not exceed a predefined threshold, avoiding unstable updates. For this model, the clipping value is set to 0.5

```
dW3 = np.clip(dW3, -gradient_clip_value, gradient_clip_value)
dW2 = np.clip(dW2, -gradient_clip_value, gradient_clip_value)
dW1 = np.clip(dW1, -gradient_clip_value, gradient_clip_value)
```

5) (Ensemble) Training Residual Updates:

Gradient clipping ensures that gradients do not exceed a predefined threshold, avoiding unstable updates. For this model, the clipping value is set to 0.5

```
def train(self, X, y, epochs=10, batch_size=32, residual_threshold=1e-3):
    residuals = y.astype(np.float64)
    for model_idx, model in enumerate(self.models):
        print(f"Training Model {model_idx + 1}/{len(self.models)}")
        model.train(X, residuals, epochs, batch_size, self.gradient_clip_value)
        predictions = model.forward(X)
        residuals -= predictions
        residual_norm = np.linalg.norm(residuals, axis=0, keepdims=True)
        if np.all(residual_norm < residual_threshold):
            print("Residuals below threshold. Stopping further training.")
            break
    residuals /= residual_norm
```

Model 3: Adam Optimizer Method

Theory

Adam optimizer leverages an adaptive learning rate optimization algorithm.

1) Initialization of Parameters:

Additional parameters for moment estimates.

```
# Initialize Adam parameters
self.m_weights = {key: np.zeros_like(value) for key, value in self.weights.items()}
self.v_weights = {key: np.zeros_like(value) for key, value in self.weights.items()}
self.m_biases = {key: np.zeros_like(value) for key, value in self.biases.items()}
self.v_biases = {key: np.zeros_like(value) for key, value in self.biases.items()}
```

2, 3) Forward Propagation, Loss Computation

Same with Model 1

Model 3: Adam Optimizer Method

4) Backward Propagation

Adopt Adam Optimization algorithm

```
def adam_update(self, gradients, params, m_params, v_params):
    self.t += 1
    updated_params = {}
    for key in params.keys():
        m_params[key] = self.beta1 * m_params[key] + (1 - self.beta1) * gradients[key]
        v_params[key] = self.beta2 * v_params[key] + (1 - self.beta2) * (gradients[key] ** 2)
        m_hat = m_params[key] / (1 - self.beta1 ** self.t)
        v_hat = v_params[key] / (1 - self.beta2 ** self.t)
        updated_params[key] = params[key] - self.learning_rate * m_hat / (np.sqrt(v_hat) + self.epsilon)
    return updated_params, m_params, v_params

def backward(self, X, y_true, y_pred):
    n_samples = y_true.shape[0]
    dZ3 = y_pred - y_true
    dW3 = np.dot(self.A2.T, dZ3) / n_samples
    db3 = np.sum(dZ3, axis=0, keepdims=True) / n_samples
    dA2 = np.dot(dZ3, self.weights["W3"].T)
    dZ2 = dA2 * self.relu_derivative(self.Z2)
    dW2 = np.dot(self.A1.T, dZ2) / n_samples
    db2 = np.sum(dZ2, axis=0, keepdims=True) / n_samples
    dA1 = np.dot(dZ2, self.weights["W2"].T)
    dZ1 = dA1 * self.relu_derivative(self.Z1)
    dW1 = np.dot(X.T, dZ1) / n_samples
    db1 = np.sum(dZ1, axis=0, keepdims=True) / n_samples
    gradients = {"W1": dW1, "W2": dW2, "W3": dW3, "b1": db1, "b2": db2, "b3": db3}
    self.weights, self.m_weights, self.v_weights = self.adam_update(
        {key: gradients[key] for key in self.weights.keys()}, self.weights, self.m_weights, self.v_weights
    )
    self.biases, self.m_biases, self.v_biases = self.adam_update(
        {key: gradients[key] for key in self.biases.keys()}, self.biases, self.m_biases, self.v_biases
    )
```

Theory

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t, \quad v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\theta_t = \theta_{t-1} - \frac{\alpha}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

Model 3: Adam Optimizer Method

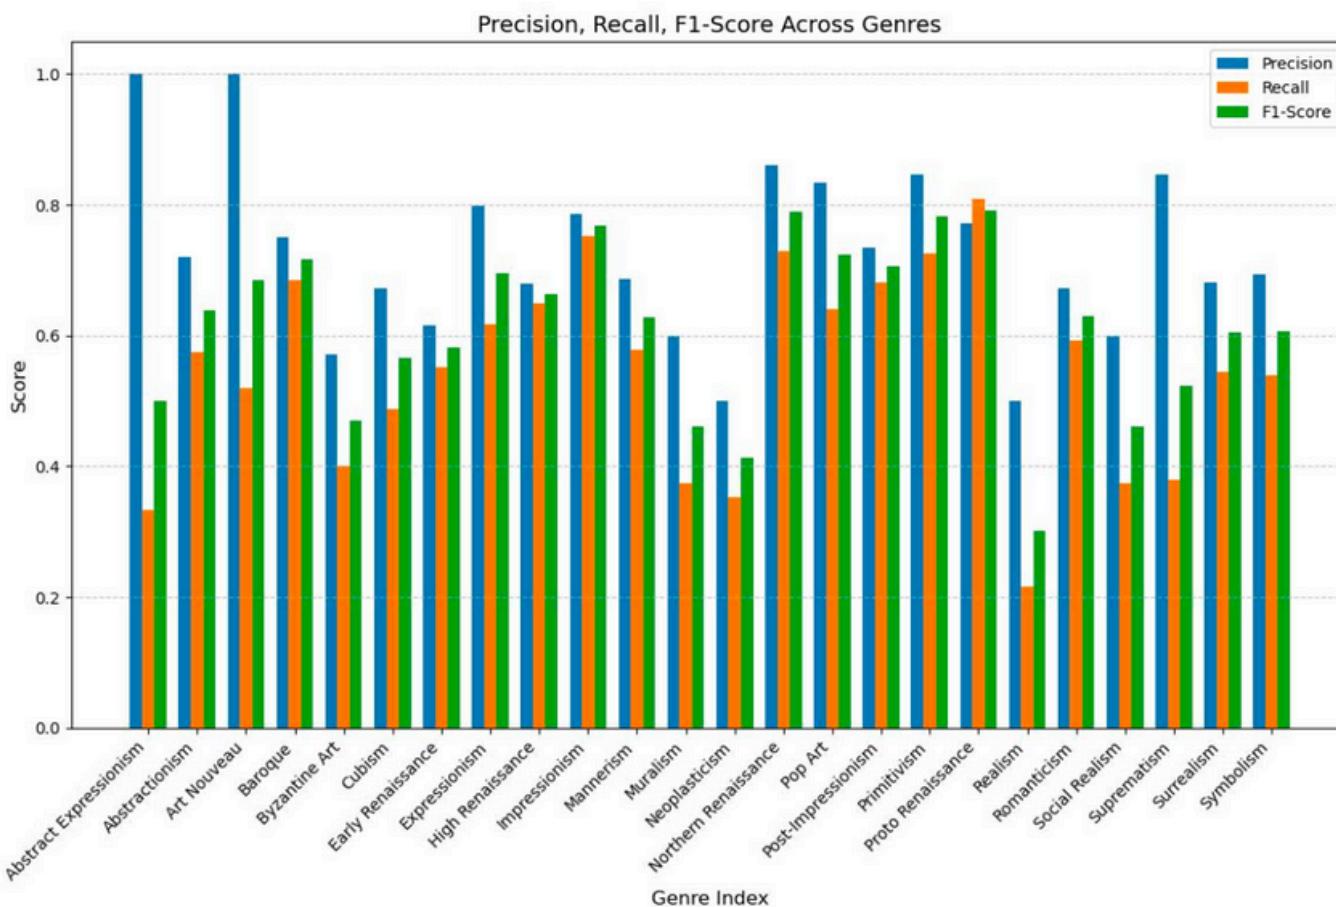
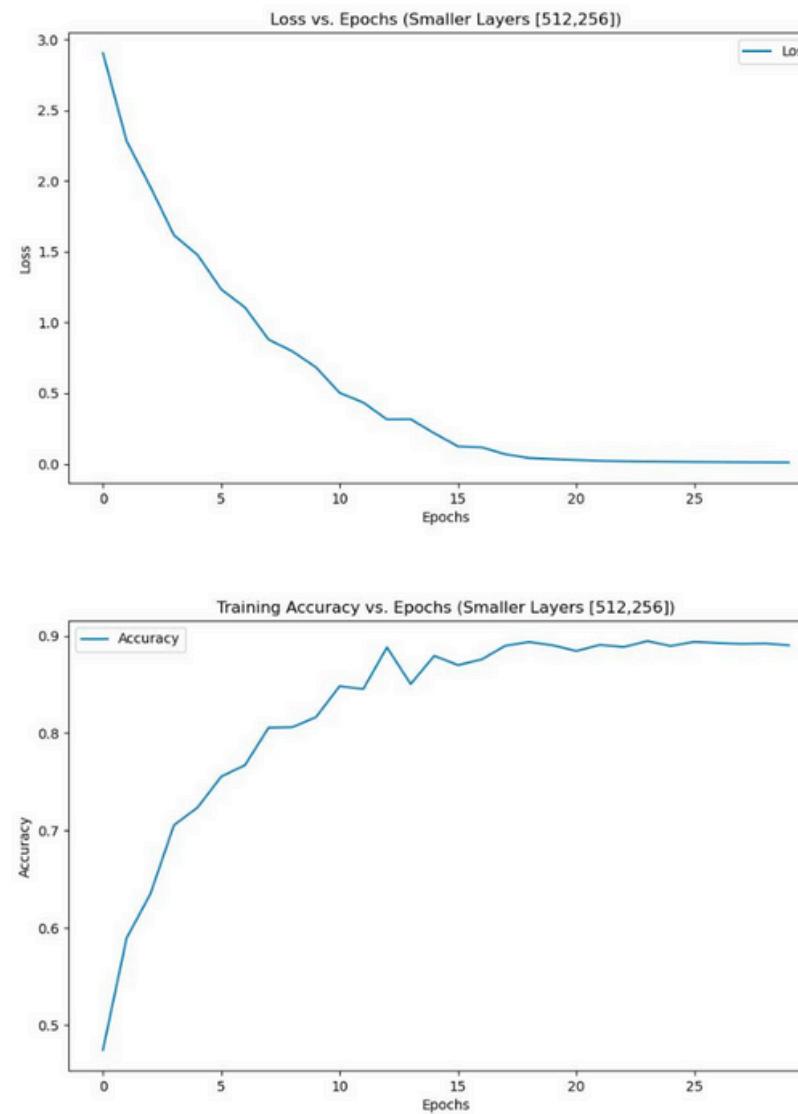
5) Training

Incorporates the unique Adam update mechanism during the parameter update.

```
def train(self, X, y, epochs=20, batch_size=32):
    n_samples = X.shape[0]
    for epoch in range(epochs):
        indices = np.arange(n_samples)
        np.random.shuffle(indices)
        X = X[indices]
        y = y[indices]
        batch_losses = []
        batch_accuracies = []
        for i in range(0, n_samples, batch_size):
            X_batch = X[i:i + batch_size]
            y_batch = y[i:i + batch_size]
            y_pred = self.forward(X_batch)
            self.backward(X_batch, y_batch, y_pred)
```

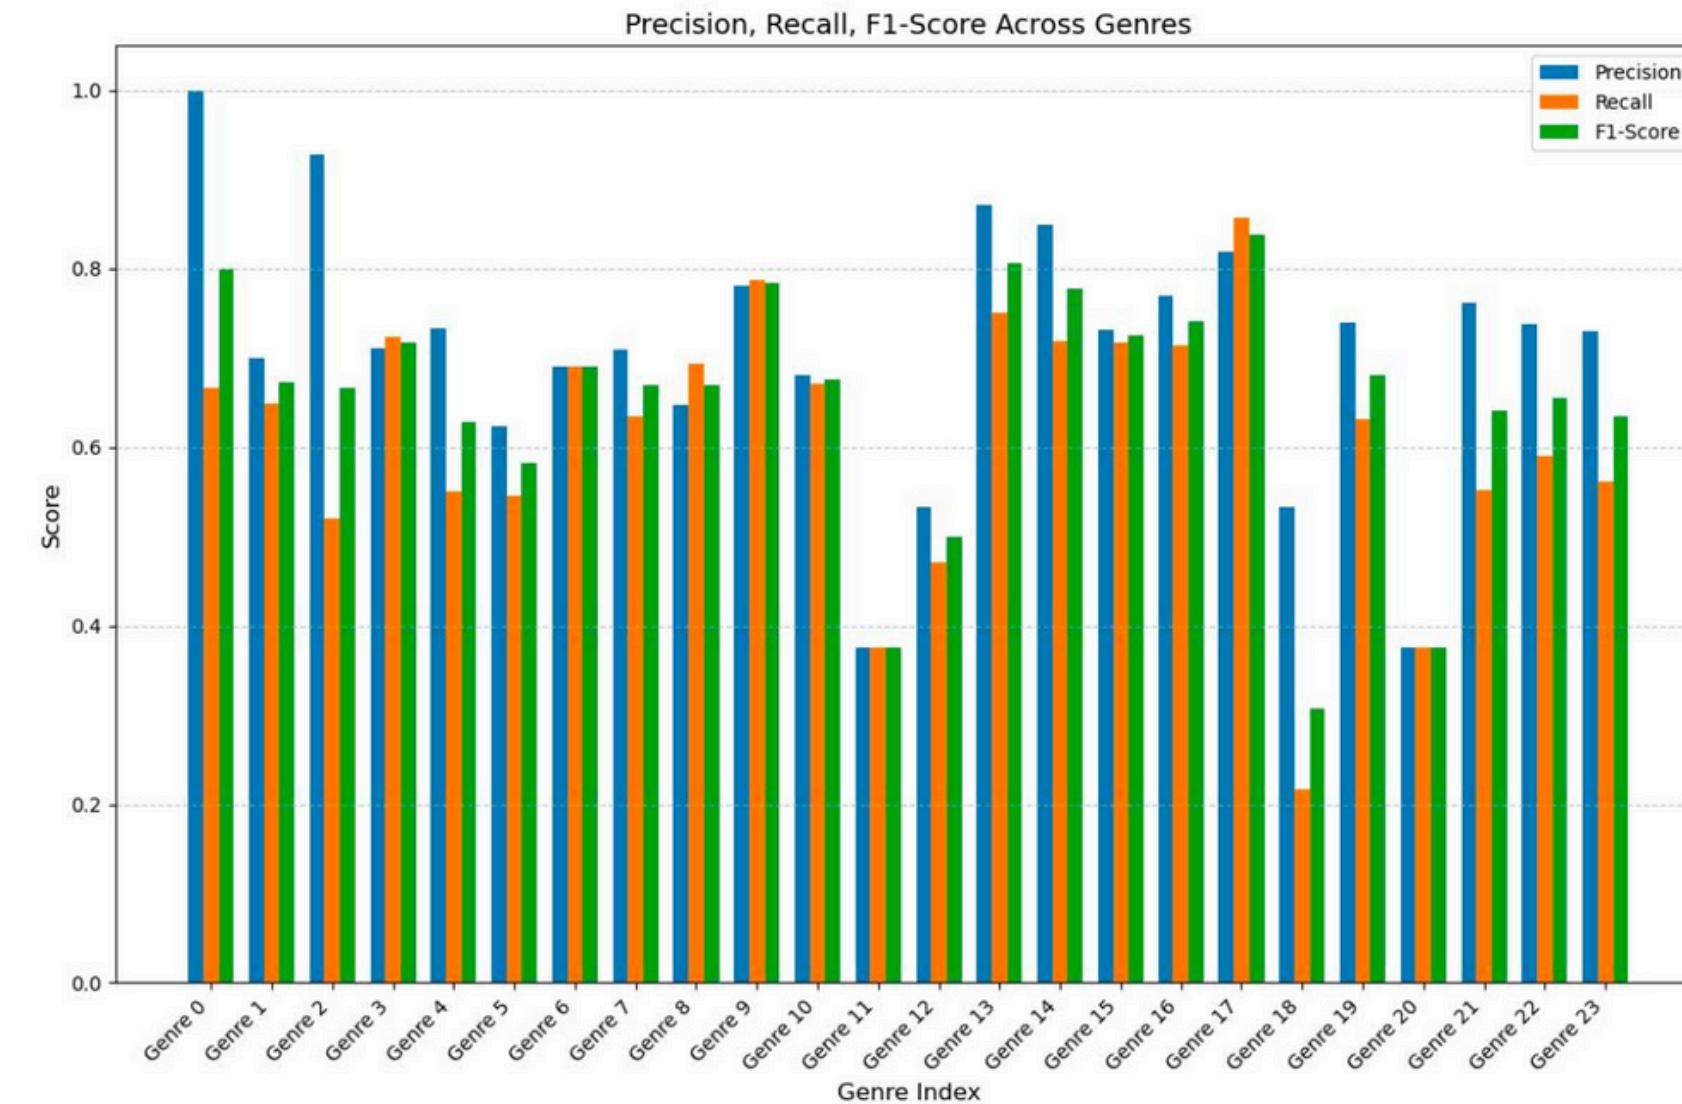
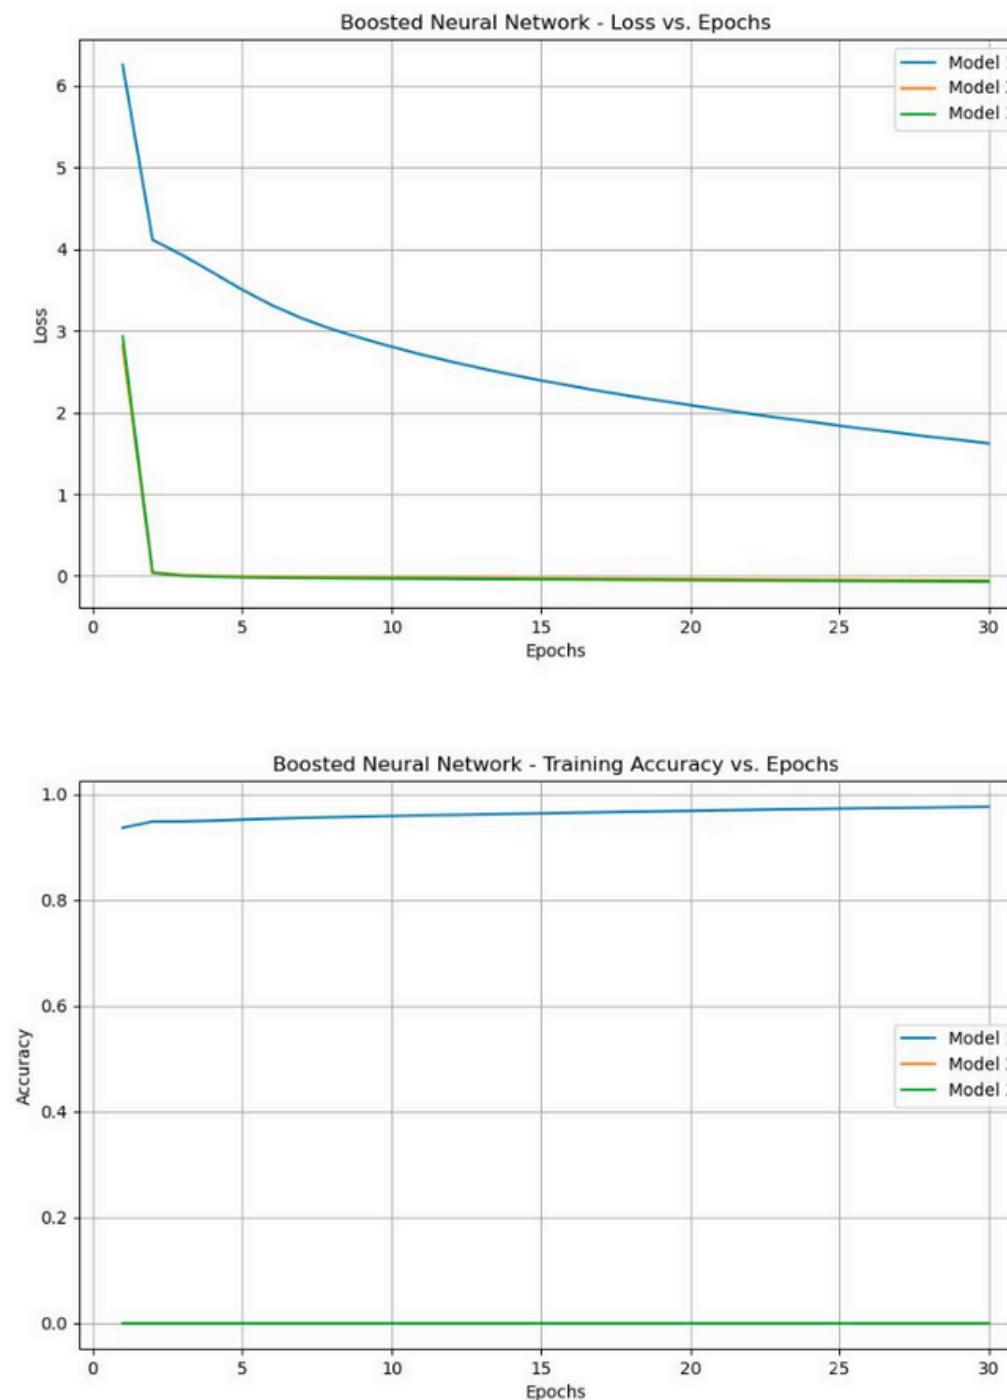
Evaluation & Results

Key results: Simple Neural Network



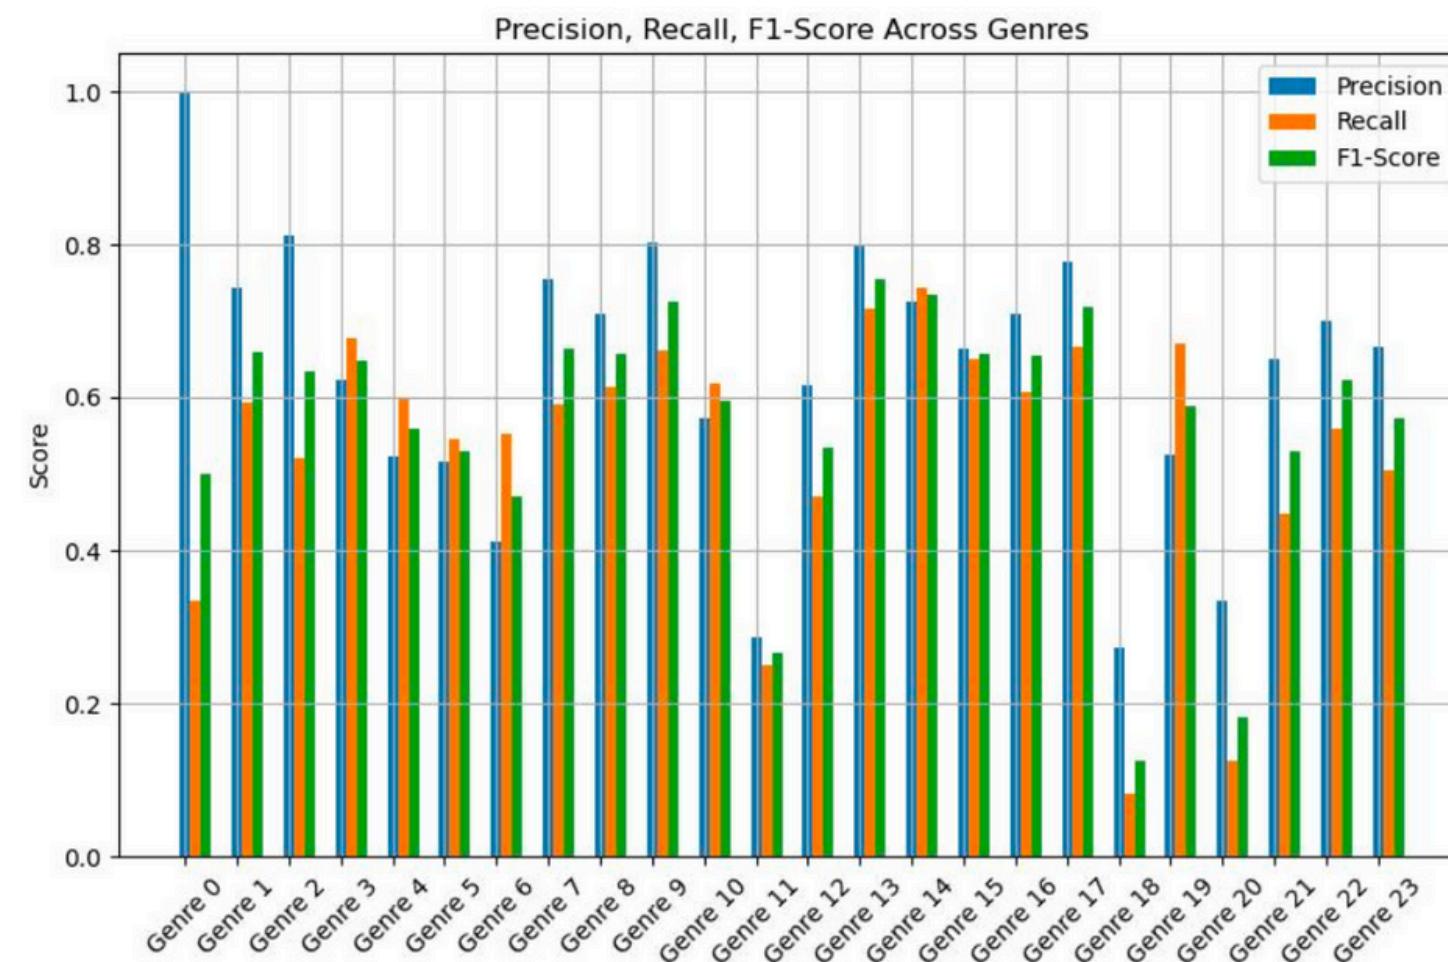
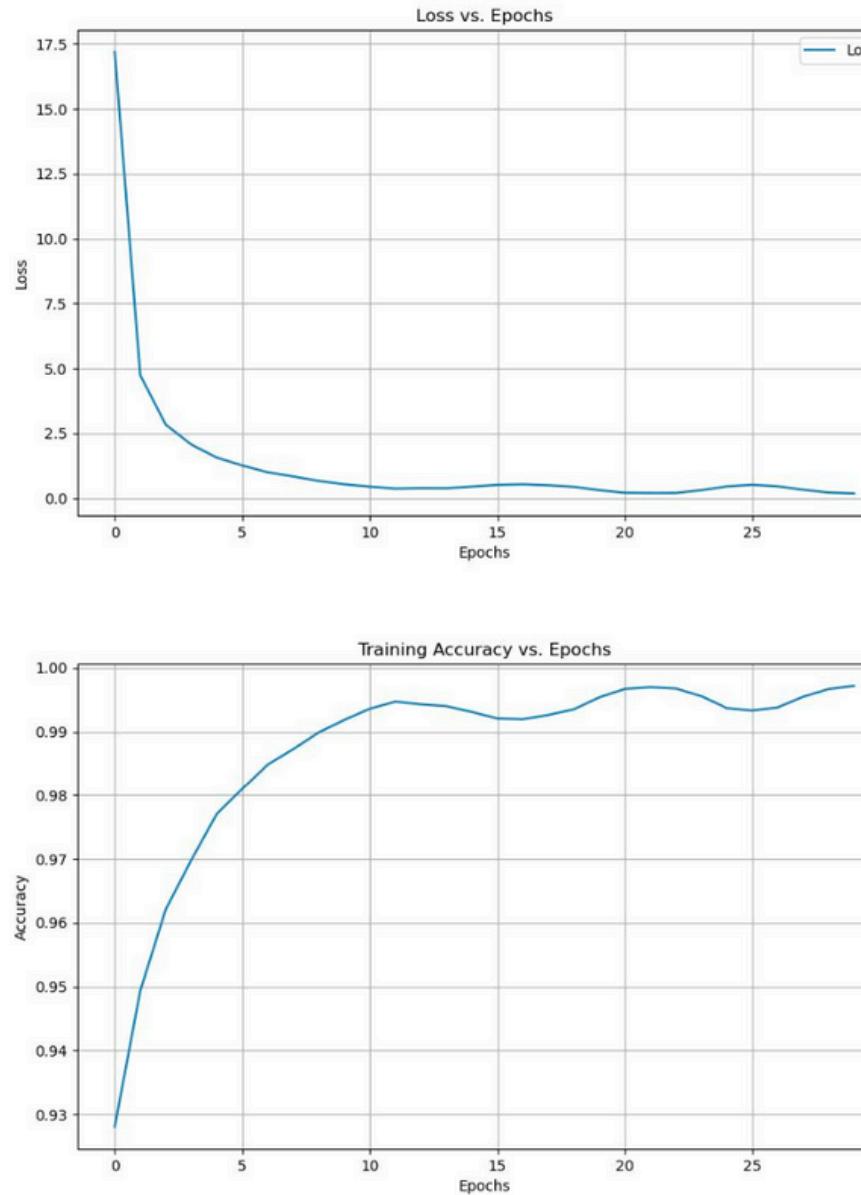
The model evaluates multi-label classification by checking if predicted genres match any associated with the artist. For the architecture [512, 256] and a learning rate of 0.01, the SNN achieved a test accuracy of 73.13%.

Key results: Boosted Neural Network



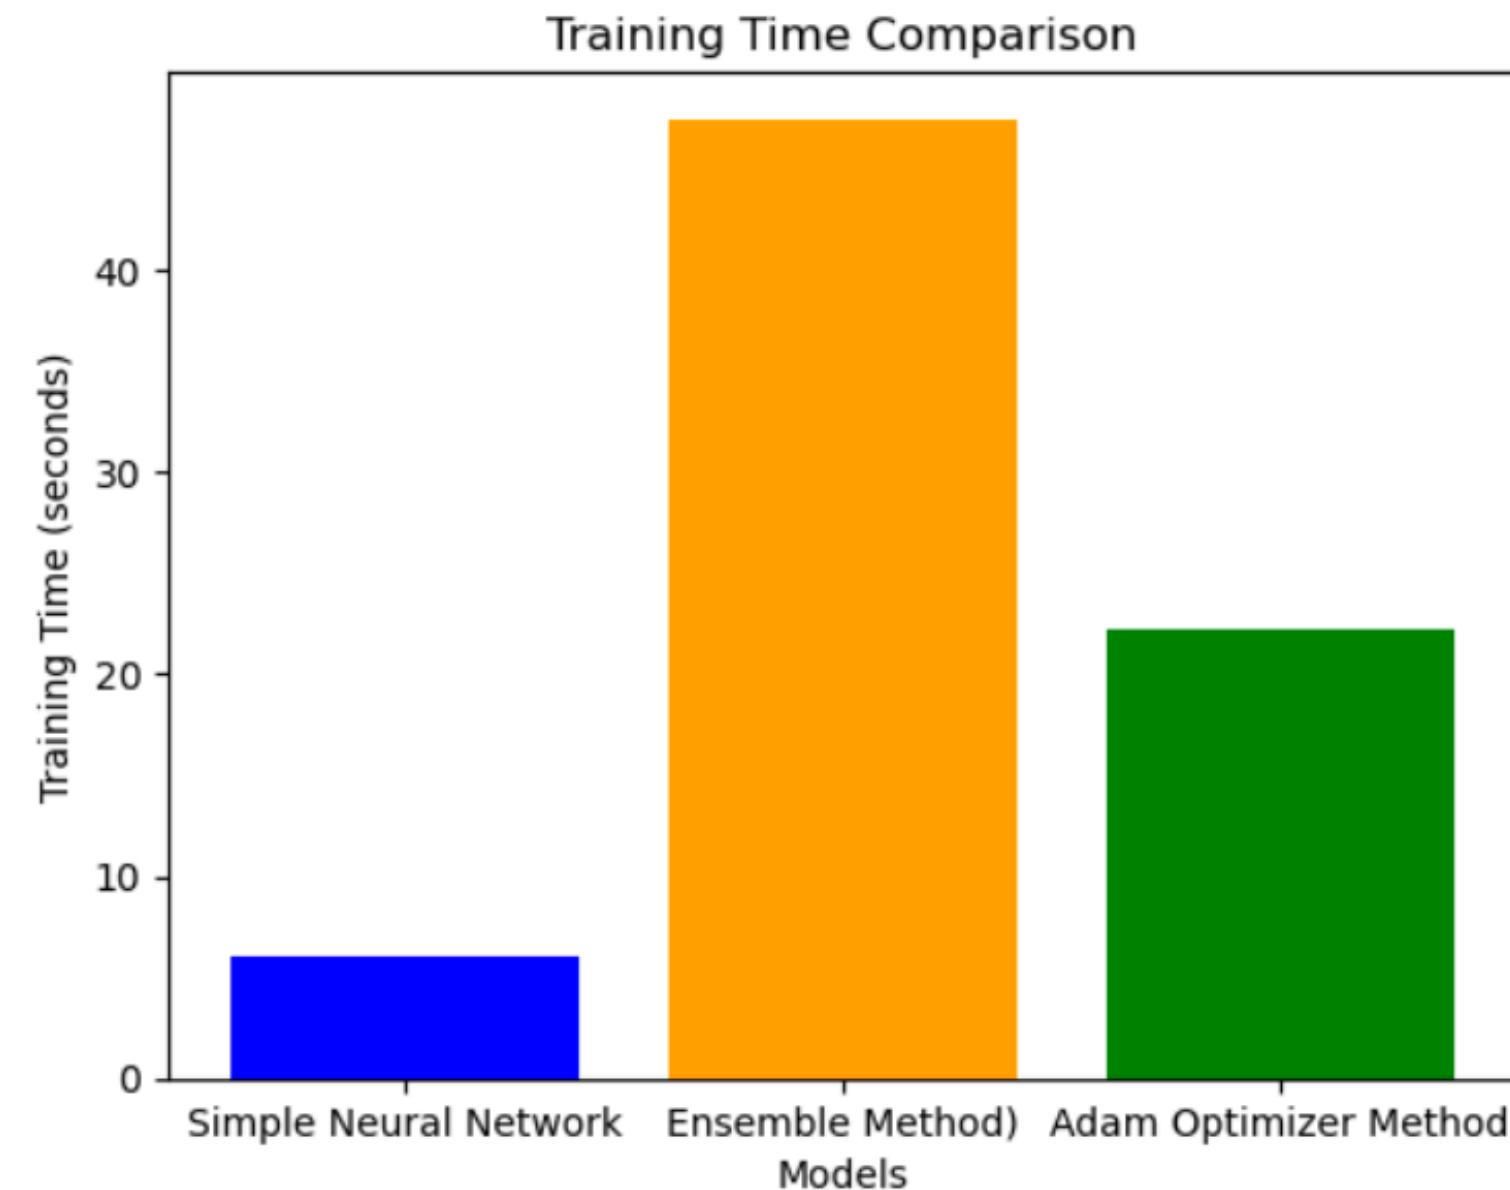
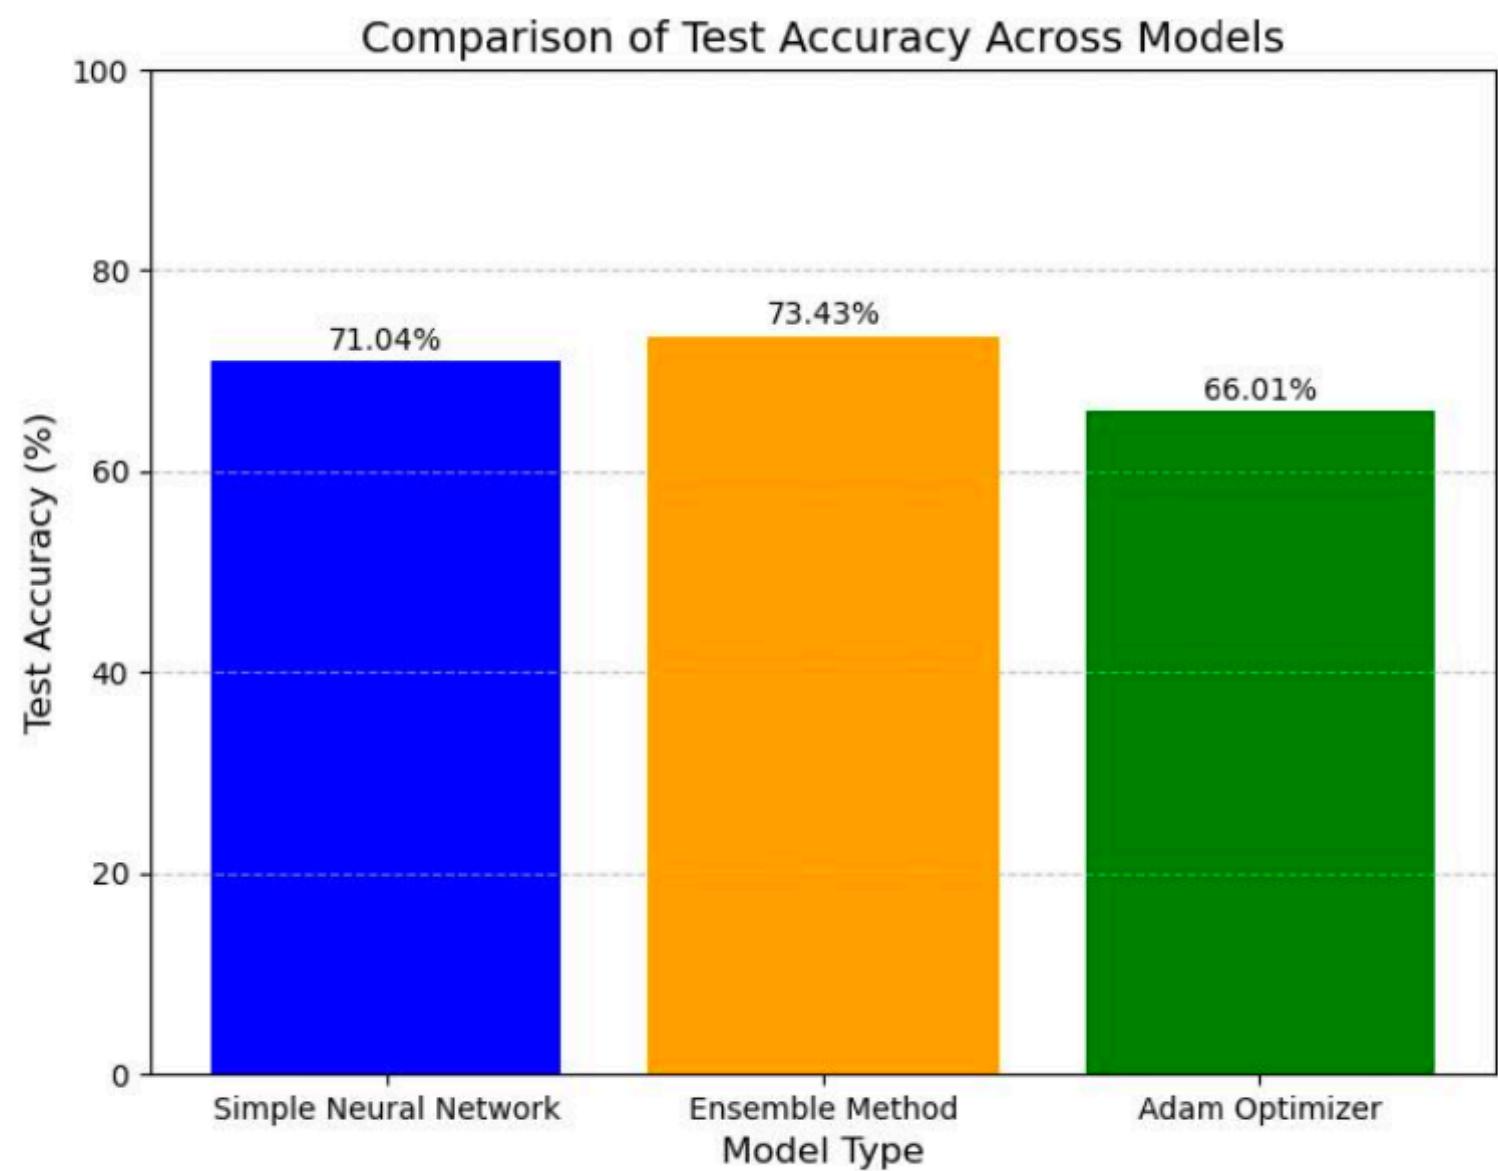
BNN achieved a test accuracy of 73.25% using the same evaluation method as SNN.

Key results: Adam Optimizer Method



Adam Optimizer Method achieved a test accuracy of 66.55% using the same evaluation method as SNN. The Adam model exhibits higher variance in precision and recall, indicating less consistent performance across genres. This reflects the Adam model's tendency to converge faster but with reduced accuracy compared to the SNN.

Comparison between all 3 models



Limitations & Potential Improvements

Limitations and improvements

Imbalanced Dataset:

Genre Imbalance Issue:

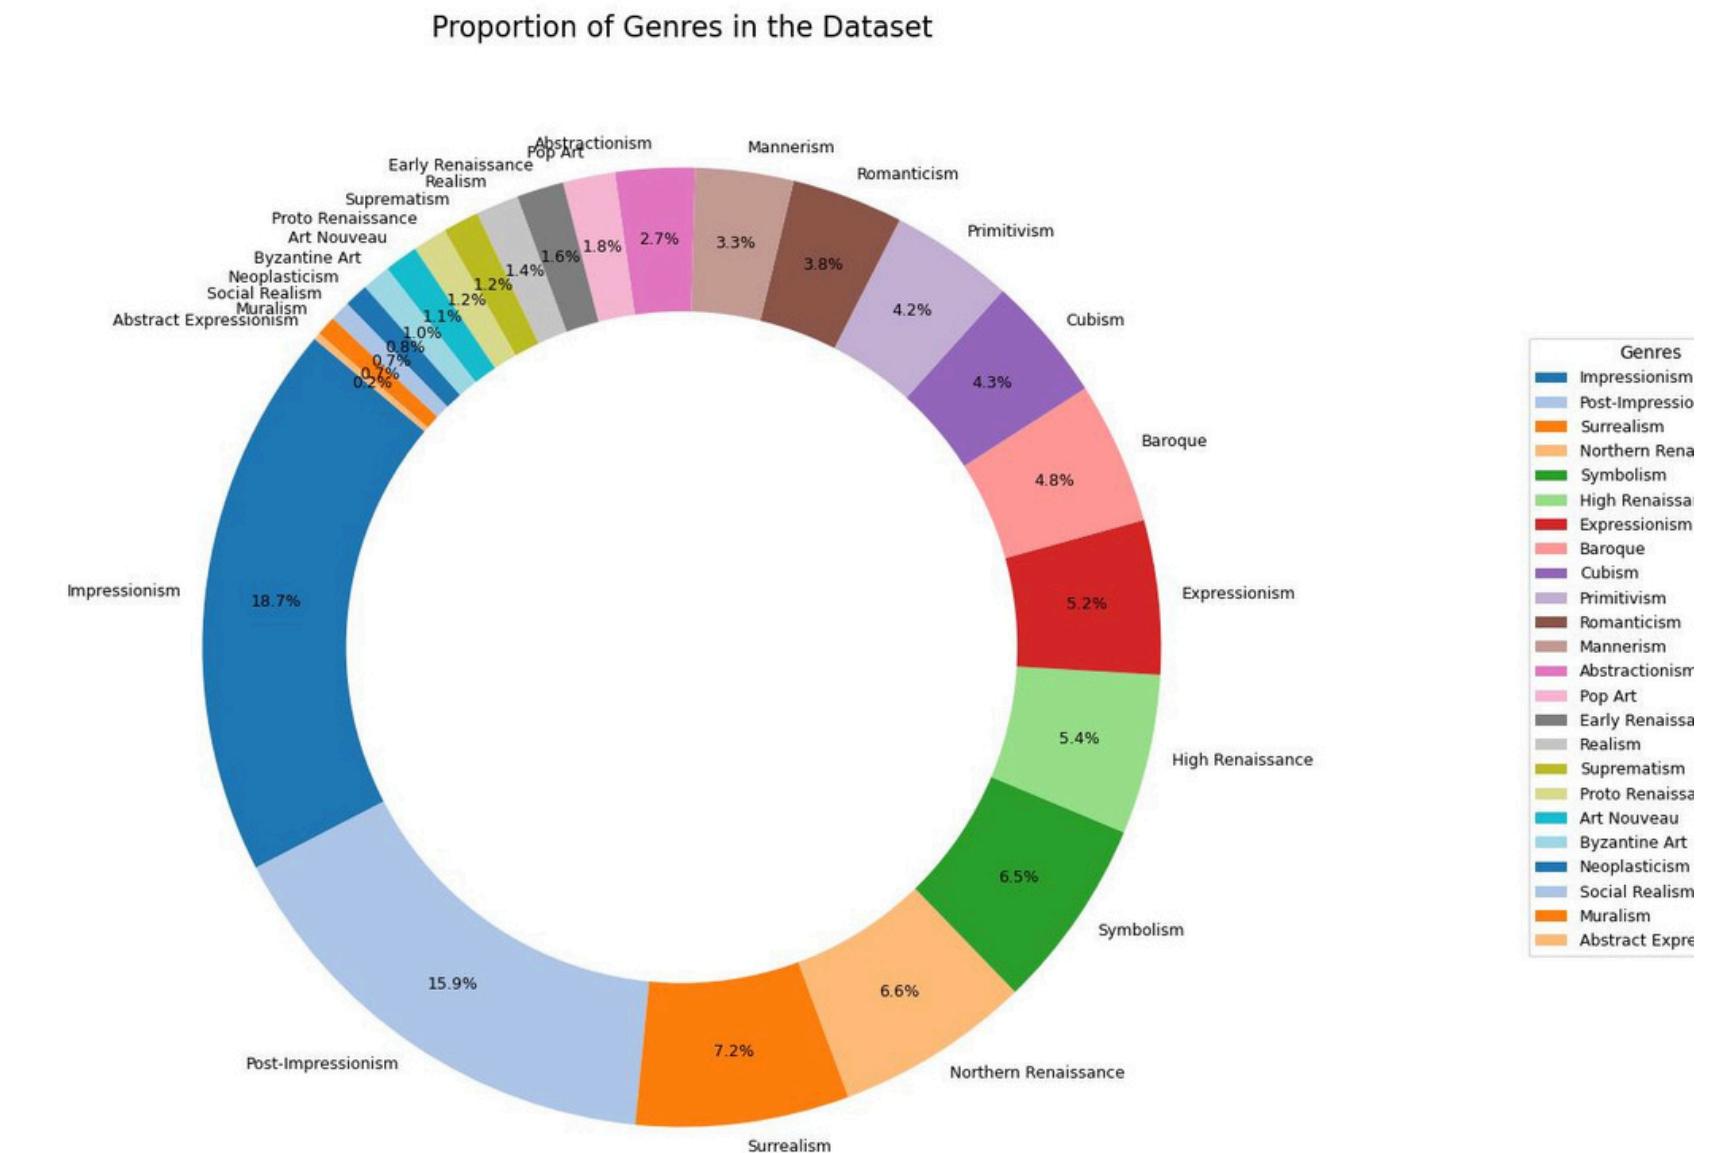
- Genres like Impressionism dominate the dataset, while others (e.g., Abstract Expressionism) are underrepresented.

Impact on Performance:

- Imbalance leads to lower Precision, Recall, and F1-Scores for minority genres.

Solutions to Improve Performance:

- Data Augmentation: Increase samples for minority genres by generating synthetic data.
- Balanced Sampling: Ensure models train on a more balanced subset of data to improve generalization across all genres.



Limitations and improvements

Imperfect Feature Extraction with ResNet50:

Pretrained ResNet50 Limitations:

- ResNet50, trained on ImageNet, excels at general object recognition.
- It may struggle to capture the nuanced stylistic details needed for art genre classification.

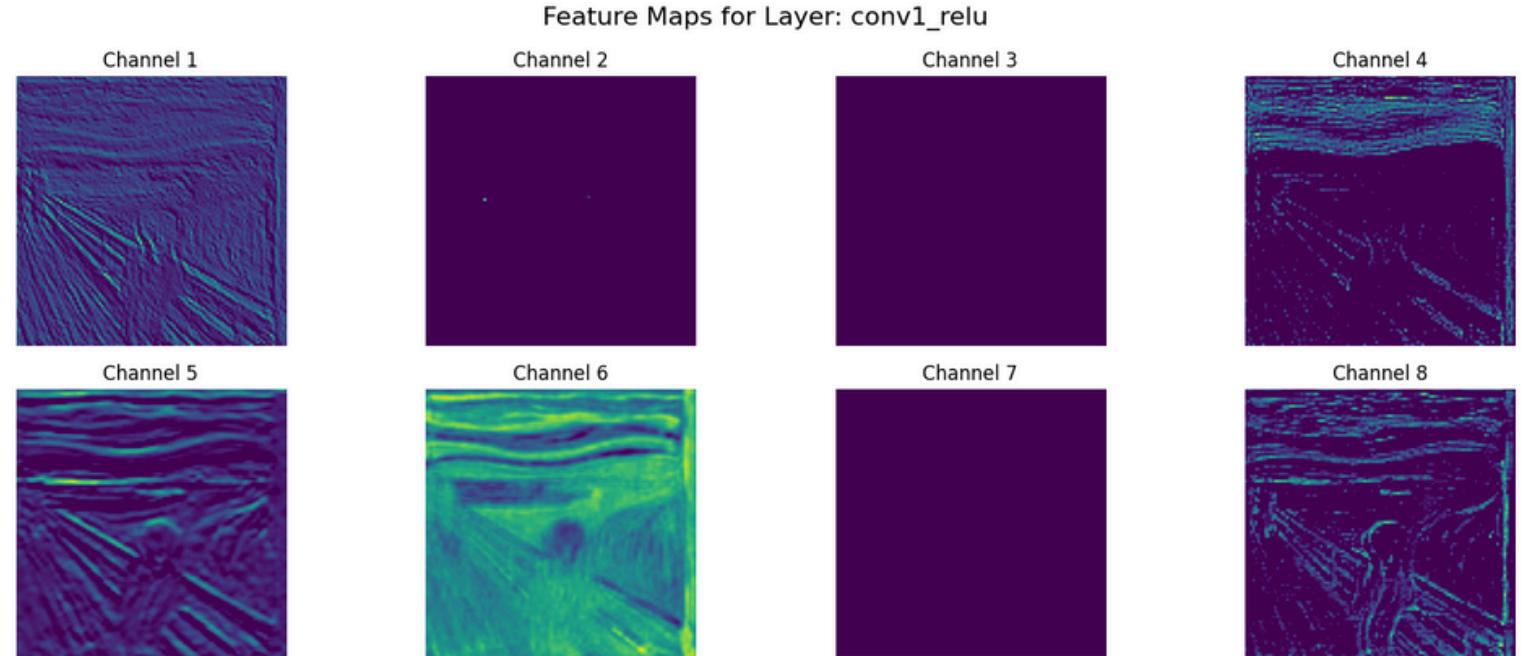


Feature Mismatch:

- Extracted features may not fully align with the specific requirements of art classification tasks.

Performance Improvement Strategies:

- Fine-Tuning: Retrain ResNet50 on domain-specific art datasets to better capture stylistic features.
- Attention Mechanisms: Integrate attention techniques to help the model focus on important stylistic details.



Conclusion

Conclusion

Research Implications:

- Demonstrates the potential of machine learning for art genre classification.
- Shows that Boosted Neural Networks (BNN) provide the best accuracy (73.43%) through ensemble learning, while Simple Neural Networks (SNN) offer efficiency and practicality.

Future Directions:

- Address Dataset Imbalance: Use balanced sampling or data augmentation to improve performance for underrepresented genres.
- Domain-Specific Fine-Tuning: Adapt ResNet50 to art datasets for better feature extraction.

Broader Impact:

- Machine learning can support art education, digital curation, and the preservation of cultural heritage by enabling efficient and accurate art classification.

